

yN_lie, sentiment, & text. R

- b.hogan@snhu.edu
- data = https://github.com/bbe2/data/blob/master/a_vs_b_data_deception.csv

R Markdown

```
library(tidytext, warn.conflicts = FALSE, quietly = TRUE)
library(stringr, warn.conflicts = FALSE, quietly = TRUE)
library(dplyr, warn.conflicts = FALSE, quietly = TRUE)
library(tidyr, warn.conflicts = FALSE, quietly = TRUE)
library(wordcloud, warn.conflicts = FALSE, quietly = TRUE)
library(ggplot2, warn.conflicts = FALSE, quietly = TRUE)
options(warn = (-1))

#####
#=> Part 0: PreProcessing + Tokenization + Stemming + Lemmatization =====
#####

hw8df0 <- read.csv(
  "C://Users//17574//Desktop//data_it304//a_vs_b_data_deception.csv",
  stringsAsFactors = FALSE)
hw8df1 <- data.frame(hw8df0)
dim(hw8df1) #92x24: ## [1] 92 24

df_text <- data.frame(1:92) #blank df
colnames(df_text) <- c("text")
as.character(df_text) ## [1] "1:92"

xtext <- as.character()
x <- 1
while (x <= 92) #put all the text & unnecessary characters into a vector
{
  yt1 <- as.character(0); yt2 <- as.character(0); yt3 <- as.character(0)
  yt4 <- as.character(0); yt5 <- as.character(0); yt6 <- as.character(0)
  yt7 <- as.character(0); yt8 <- as.character(0); yt9 <- as.character(0)
  yt10 <- as.character(0); yt11 <- as.character(0); yt12 <- as.character(0)
  yt13 <- as.character(0); yt14 <- as.character(0); yt15 <- as.character(0)
  yt16 <- as.character(0); yt17 <- as.character(0); yt18 <- as.character(0)
  yt19 <- as.character(0); yt20 <- as.character(0); yt21 <- as.character(0)
  yt22 <- as.character(0)
  yt1 <- hw8df1[x+1,3]; yt2 <- hw8df1[x+1,4]; yt3 <- hw8df1[x+1,5]
  yt4 <- hw8df1[x+1,6]; yt5 <- hw8df1[x+1,7]; yt6 <- hw8df1[x+1,8]
  yt7 <- hw8df1[x+1,9]; yt8 <- hw8df1[x+1,10]; yt9 <- hw8df1[x+1,11]
  yt10 <- hw8df1[x+1,12]; yt11 <- hw8df1[x+1,13]; yt12 <- hw8df1[x+1,14]
  yt13 <- hw8df1[x+1,15]; yt14 <- hw8df1[x+1,16]; yt15 <- hw8df1[x+1,17]
  yt16 <- hw8df1[x+1,18]; yt17 <- hw8df1[x+1,19]; yt18 <- hw8df1[x+1,20]
  yt19 <- hw8df1[x+1,21]; yt20 <- hw8df1[x+1,22]; yt21 <- hw8df1[x+1,23]
  yt22 <- hw8df1[x+1,24]
  xtext <- as.character(xtext)
  xtext <- paste(yt1,yt2,yt3,yt4,yt5,yt6,yt7,yt8,yt9,yt10,yt11,yt12,yt13,
    yt14,yt15,yt16,yt17,yt18,yt19,yt20,yt21,yt22)
  df_text[x,1] <- xtext
  x <- x+1 }
}
```

```

#dataframe with lie, sentiment, text
list1 <- c(1:nrow(df_text)) #create analysis numeric ID
lie_value <- rep(99, length(list1))
sent_value <- rep(99, length(list1))
reviewID <- rep(1:length(list1))
df_text <- cbind(hw8df1[,1],hw8df1[,2],df_text,reviewID,lie_value,
                sent_value) #add lie & sentiment back
colnames(df_text) <- c("lie","sentiment","text","reviewID", "lievalue","sentv
alue")
remove(list1, lie_value, sent_value, reviewID, x)
i <- 1
while (i <=length(df_text)) #put all text & un-nessary chrs in vector
{ if (df_text[i,1]=="t") {df_text[i,5]=1}
  if (df_text[i,1]=="f") {df_text[i,5]=0}
  if (df_text[i,2]=="n") {df_text[i,6]=0}
  if (df_text[i,2]=="p") {df_text[i,6]=1}
  i <- i+1 }

#=====> Part 1: STEMMING, LEMMATIZATION, AND FREQUENCY INSPECTION
#=====

review_words <- df_text %>% #==> Tokenization
  select(lie, sentiment, text, reviewID) %>%
  unnest_tokens(word, text, to_lower=TRUE) %>%
  count(lie, sentiment, reviewID, word, sort=FALSE) %>%
  bind_tf_idf(word,reviewID,n)
# spread(key=word, value=tf_idf)
head(review_words,7) #review of 10 words from tokens

## # A tibble: 7 x 8
##   lie sentiment reviewID word          n      tf      idf  tf_idf
##   <fct> <fct>      <int> <chr>      <int> <dbl> <dbl> <dbl>
## 1 f      n          1 a          2 0.0328 0.280  0.00919
## 2 f      n          1 also        3 0.0492 2.42   0.119
## 3 f      n          1 american    1 0.0164 3.11   0.0510
## 4 f      n          1 and         2 0.0328 0.0810 0.00265
## 5 f      n          1 are         1 0.0164 1.50   0.0247
## 6 f      n          1 area        1 0.0164 3.11   0.0510
## 7 f      n          1 back        1 0.0164 1.86   0.0305

#==>STEMMING & LEMMATIZATION
library(SnowballC)
review_words$word <- wordStem(review_words$word)
#==> REMOVE STOP WORDS ==> (6991-2941)= 4050 words removed
nrow(review_words) #6991

## [1] 4833

review_words <- review_words %>%
  filter(!word %in% stop_words$word, str_detect(word,"^[a-z']+$"))
nrow(review_words) #2941 <high-ish>

## [1] 2448

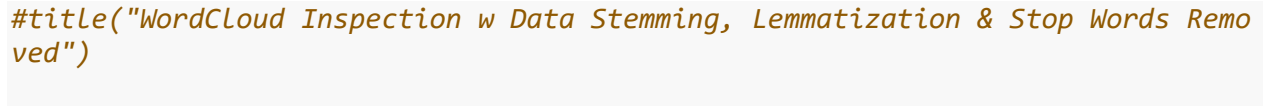
```

```
#==> VISUALIZATION frequency
review_words %>%
  count(word) %>%
  with(wordcloud(word, n, max.words=200)) #well that is pretty darn cool
```

```
review_words %>%
  count(word) %>%
  with(wordcloud(word, n, max.words=200)) #well that is pretty darn cool
```

```
review_words %>%
  count(word) %>%
  with(wordcloud(word, n, max.words=200)) #well that is pretty darn cool
```

```
review_words %>%
  count(word) %>%
  with(wordcloud(word, n, max.words=200)) #well that is pretty darn cool
```



```
#title("WordCloud Inspection w Data Stemming, Lemmatization & Stop Words Remo  
ved")
```

```

#=====
#=>Chapter 2:TidyText - CAST_DTM=> Tidy=> Cast-Sparse=> Matrix=> DataFrame !
#=====
library(tidytext,warn.conflicts = FALSE, quietly = TRUE)
review_words$reviewID <- as.character(review_words$reviewID)
mydata <- data.frame(review_words)
str(mydata)

## 'data.frame': 2448 obs. of 8 variables:
## $ lie : Factor w/ 2 levels "f","t": 1 1 1 1 1 1 1 1 1 1 ...
## $ sentiment: Factor w/ 2 levels "n","p": 1 1 1 1 1 1 1 1 1 1 ...
## $ reviewID : chr "1" "1" "1" "1" ...
## $ word : chr "american" "ar" "buffet" "cheap" ...
## $ n : int 1 1 1 1 1 1 1 1 1 1 ...
## $ tf : num 0.0164 0.0164 0.0164 0.0164 0.0164 ...
## $ idf : num 3.11 1.5 3.4 3.4 2.89 ...
## $ tf_idf : num 0.051 0.0247 0.0558 0.0558 0.0474 ...

mydtm <- cast_dtm(mydata, reviewID, word, n) #row is doc, col are words
df_tidy <- tidy(mydtm) #create df and tidy
df_cast <- cast_sparse(df_tidy,document,term,count) # create matrix
dfx <- as.matrix(df_cast)
dfxx <- data.frame(dfx) # word cube for ml
#str(dfxx)
#=====
#=>Chapter 3: Machine Learn PreProcess: merge lie/sentim. Create Datasets
#=====
library(caret,warn.conflicts = FALSE, quietly = TRUE) #multiple...algorithms
set.seed(199)
#use row names to merge back lie & sentiment lables
dfxx <- cbind(reviewID = rownames(dfxx),dfxx) #fix & sort row names
dfxx$reviewID <- as.numeric(dfxx$reviewID) #make back numeric
df_labels <- data.frame(df_text[,c(1,4)]) #1=lie, 2=sentiment, 4=reviewID
df_lie <-merge(dfxx,df_labels,key="reviewID") #only merge target
df_labels <- data.frame(df_text[,c(2,4)]) #1=lie, 2=sentiment, 4=reviewID
df_sentiment <-merge(dfxx,df_labels,key="reviewID") #lie
labels_ID <-df_lie[,1] #need to remove ID for running models
df_lie <- df_lie[,c(-1)] #now remove ID as data all merged
df_sentiment <- df_sentiment[,c(-1)] #now remove ID as data all merged
write.csv(df_lie,"today.csv") #makes sure no duplicates on merge/funny
#Create the datasets
train_index <- createDataPartition(df_lie$lie, p=0.7, list=FALSE)
df_lie_train <-df_lie[train_index,]
df_lie_test <- df_lie[-train_index,] #split data set training...
df_lie_test_labels <-df_lie_test$lie
df_lie_test$sentiment <- as.factor(c("?"))
df_sentiment_train <-df_sentiment[train_index,]
df_sentiment_test <- df_sentiment[-train_index,] #split data set training...
df_sentiment_test_labels <-df_sentiment_test$sentiment
df_sentiment_test$sentiment <- as.factor(c("?"))
#str(df_sentiment_test);dim(df_sentiment_test)

```

```

#=====
#=> Multinomial NB Analysis =====
#=====
library(klaR, warn.conflicts = FALSE, quietly = TRUE)
## Attaching package: 'MASS'

## The following object is masked from 'package:dplyr':
##
##      select

mlie_nb1 <- train(lie~., data=df_lie_train, method="nb",
                  trControl = trainControl(method="cv", number=3),
                  tuneGrid=expand.grid(fL=1:3, usekernel=c(TRUE, FALSE), adjust=1:3))
                                #Laplace=fL; usekernel=smoothing
mlie_predict_nb1 <- predict(mlie_nb1, newdata=df_lie_test, type="raw")
confusionMatrix(mlie_predict_nb1, df_lie_test$lie)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  f  t
##           f 13 13
##           t  0  0
##
##              Accuracy : 0.5                      =====> yikes!
##              95% CI : (0.2993, 0.7007)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : 0.5774905
##
##              Kappa : 0
##
##      McNemar's Test P-Value : 0.0008741
##
##              Sensitivity : 1.0
##              Specificity : 0.0
##              Pos Pred Value : 0.5
##              Neg Pred Value : NaN
##              Prevalence : 0.5
##              Detection Rate : 0.5
##      Detection Prevalence : 1.0
##              Balanced Accuracy : 0.5
##
##              'Positive' Class : f

#sentiment
msentiment_nb1 <- train(sentiment~., data=df_sentiment_train, method="nb",
                        trControl = trainControl(method="cv", number=3),
                        tuneGrid=expand.grid(fL=1:3, usekernel=c(TRUE, FALSE), adjust
=1:3))
msentiment_predict_nb1 <- predict(msentiment_nb1, newdata=df_sentiment_test, t
ype="raw")

```

```

#=====
#==> SVM: Preproces(Normalize)
#=====
#normalize and get the lie target variable back in dataframes
normalize <- function(x) { (x - mean(x))/sd(x) } #z-score transformation
ndf_lie_train <- as.data.frame(lapply(df_lie_train[,1:1013],normalize)) #remove
ve row 1
ndf_lie_train[is.na(ndf_lie_train)] <-0 #remove NaN resulting in rows w zero
x <- data.frame(df_lie_train[,1014])
colnames(x) <- c("lie")
ndf_lie_train <- data.frame(ndf_lie_train,x)#normalize lie test data frames
ndf_lie_test <- as.data.frame(lapply(df_lie_test[,1:1013],normalize))
ndf_lie_test[is.na(ndf_lie_test)] <-0 #remove NaN resulting in rows w zeros
x <- data.frame(df_lie_test[,1014])
colnames(x) <- c("lie")
ndf_lie_test <- data.frame(ndf_lie_test,x)
#write.csv(ndf_sentiment_test,"today.csv")
#=====>sentiment
ndf_sentiment_train <- as.data.frame(lapply(df_sentiment_train[,1:1013],norma
lize)) #remove row 1
ndf_sentiment_train[is.na(ndf_sentiment_train)] <-0 #remove NaN resulting in
rows
x <- data.frame(df_sentiment_train[,1014])
colnames(x) <- c("sentiment")
ndf_sentiment_train <- data.frame(ndf_sentiment_train,x)#normalize test df
ndf_sentiment_test <- as.data.frame(lapply(df_sentiment_test[,1:1013],normali
ze))
ndf_sentiment_test[is.na(ndf_sentiment_test)] <-0 #remove NaN in rows w zero
x <- data.frame(df_sentiment_test[,1014])
colnames(x) <- c("sentiment")
ndf_sentiment_test <- data.frame(ndf_sentiment_test,x)
remove(x) #write.csv(ndf_sentiment_test,"today.csv")

#=====
#==> SVM: LIE Models
#=====
=
library(gmodels,warn.conflicts = FALSE, quietly = TRUE) #for chi-square
set.seed(1984)
mlie_svm <-train(lie ~., data=ndf_lie_train, method= "svmLinear",
                 tuneGrid = expand.grid(C=seq(0,1,0.05)),
                 trControl = trainControl(method = "boot",
                                           number=10))

#mlie_svm #==> C (cost)
mlie_predict_svm_linear <- predict(mlie_svm, newdata=ndf_lie_test)
confusionMatrix(mlie_predict_svm_linear, ndf_lie_test$lie)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction f t
##           f 5 7

```

```

##          t 8 6
##
##          Accuracy : 0.4231
##          95% CI : (0.2335, 0.6308)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : 0.8365
##
##          Kappa : -0.1538
##
##      McNemar's Test P-Value : 1.0000
##
##          Sensitivity : 0.3846
##          Specificity : 0.4615
##          Pos Pred Value : 0.4167
##          Neg Pred Value : 0.4286
##          Prevalence : 0.5000
##          Detection Rate : 0.1923
##      Detection Prevalence : 0.4615
##          Balanced Accuracy : 0.4231
##
##      'Positive' Class : f
##
#chi-square
CrossTable(mlie_predict_svm_linear,df_lie_test_labels,
           prop.chisq = FALSE, prop.t = FALSE, dnn=c('predicted','actual'))
##
##      Cell Contents
##      |-----|
##      |                N |
##      |      N / Row Total |
##      |      N / Col Total |
##      |-----|
##
##      Total Observations in Table:  26
##
##
##      predicted | actual      f |          t | Row Total |
##      -----|-----|-----|-----|
##              f |          5 |          7 |         12 |
##              | 0.417 | 0.583 | 0.462 |
##              | 0.385 | 0.538 |      |
##      -----|-----|-----|-----|
##              t |          8 |          6 |         14 |
##              | 0.571 | 0.429 | 0.538 |
##              | 0.615 | 0.462 |      |
##      -----|-----|-----|-----|
##      Column Total |          13 |          13 |         26 |
##              | 0.500 | 0.500 |      |
##      -----|-----|-----|-----|

```

```

#preceision & recall
# precision_mlie_svm <-posPredValue(mlie_predict_svm_linear,ndf_lie_test$lie,
positive="yes")
# recall_mlie_svm <-sensitivity(mlie_predict_svm_linear,ndf_lie_test$lie,positi
tive="yes")
# precision_mlie_svm;recall_mlie_svm

#==> SVM w non-linear Kernel RBF
mlie_svm_rbf <- train(lie ~., data=ndf_lie_train, method= "svmRadial",
                      tuneGrid = expand.grid(sigma=seq(0,1,0.1),
                                              C=seq(0,1,0.1)),
                      trControl = trainControl(method = "boot",
                                              number=10))

#mlie_svm_rbf
lie_predict_svm_rbf <- predict(mlie_svm_rbf, newdata=ndf_lie_test)
#chi-square
CrossTable(lie_predict_svm_rbf,df_lie_test_labels,
           prop.chisq = FALSE, prop.t = FALSE, dnn=c('predicted','actual'))

##
##
##      Cell Contents
## |-----|
## |                      N |
## |-----|
##
##
## Total Observations in Table:  26
##
##
##      | df_lie_test_labels
## lie_predict_svm_rbf |      f      |      t      | Row Total |
## -----|-----|-----|-----|
##                      f |      13      |      13      |      26      |
## -----|-----|-----|-----|
##      Column Total |      13      |      13      |      26      |
## -----|-----|-----|-----|
##

```



```

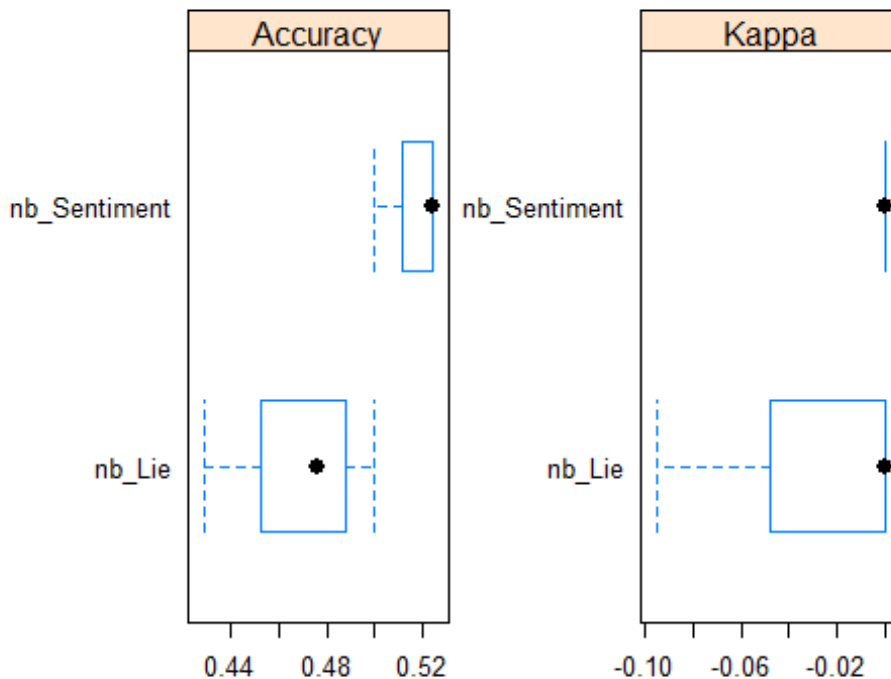
#=====
#==>  SVM: SENTIMENT MODELS
#=====
set.seed(1984)
m_sentiment_svm <- train(sentiment ~., data=ndf_sentiment_train, method= "svmL
linear",
                        tuneGrid = expand.grid(C=seq(0,1,0.05)),
                        trControl = trainControl(method = "boot",
                                                number=10))
#m_sentiment_svm                                #==>  C (cost)
m_sentiment_predict_svm_linear <- predict(m_sentiment_svm, newdata=ndf_sentim
ent_test)
#Chi-square
CrossTable(m_sentiment_predict_svm_linear,df_sentiment_test_labels,
           prop.chisq = FALSE, prop.t = FALSE, dnn=c('predicted','actual'))

#=====
#==>  MODEL COMPARISON GRAPHS
#=====
m_nb_compare <- resamples(list(nb_Lie=mlie_nb1, nb_Sentiment=msentiment_nb1))
summary(m_nb_compare)

##
## Call:
## summary.resamples(object = m_nb_compare)
##
## Models: nb_Lie, nb_Sentiment
## Number of resamples: 3
##
## Accuracy
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## nb_Lie      0.4285714 0.4523810 0.4761905 0.468254 0.4880952 0.5000000
## nb_Sentiment 0.5000000 0.5119048 0.5238095 0.515873 0.5238095 0.5238095
##           NA's
## nb_Lie      0
## nb_Sentiment 0
##
## Kappa
##           Min.   1st Qu. Median     Mean 3rd Qu. Max. NA's
## nb_Lie      -0.09565217 -0.04782609      0 -0.03188406      0      0      0
## nb_Sentiment 0.00000000 0.00000000      0 0.00000000      0      0      0

scales <- list(x=list(relation = "free"),
              y=list(relation = "free"))
bwplot(m_nb_compare, scales = scales)

```



```
m_SVM_compare <- resamples(list(svm_Lie=mlie_svm, svm_RBF_Lie=mlie_svm_rbf,
                                svm_Sentiment=m_sentiment_svm, svm_RBF_Sentiment=m_sentiment_svm_rbf
                                f ))
summary(m_SVM_compare)
```

```
##
## Call:
## summary.resamples(object = m_SVM_compare)
##
## Models: svm_Lie, svm_RBF_Lie, svm_Sentiment, svm_RBF_Sentiment
## Number of resamples: 10
##
## Accuracy
##           Min.   1st Qu.   Median     Mean   3rd Qu.
## svm_Lie      0.3043478 0.4498433 0.4782609 0.4491517 0.4950000
## svm_RBF_Lie   0.3478261 0.4360870 0.5000000 0.4824461 0.5138889
## svm_Sentiment 0.4333333 0.4460870 0.5086207 0.5197373 0.5395257
## svm_RBF_Sentiment 0.3478261 0.3934783 0.4120370 0.4401471 0.4524457
##           Max. NA's
## svm_Lie      0.5000000 0
## svm_RBF_Lie   0.5833333 0
## svm_Sentiment 0.7391304 0
## svm_RBF_Sentiment 0.6153846 0
##
## Kappa
##           Min.   1st Qu.   Median     Mean
## svm_Lie      -0.1774744 -0.05376623 0.003759398 -0.007211927
## svm_RBF_Lie   0.0000000 0.00000000 0.000000000 0.009448819
## svm_Sentiment -0.1767442 0.04543807 0.056782334 0.087421632
```

```
## svm_RBF_Sentiment 0.0000000 0.00000000 0.000000000 0.000000000
##                               3rd Qu.      Max. NA's
## svm_Lie            0.06515499 0.08906883    0
## svm_RBF_Lie        0.00000000 0.09448819    0
## svm_Sentiment       0.13380404 0.42975207    0
## svm_RBF_Sentiment 0.00000000 0.00000000    0

scales <- list(x=list(relation = "free"),
               y=list(relation = "free"))
bwplot(m_SVM_compare, scales = scales)
```

